

# Neural Networks Part 2

Jim Rehg

Based on slides prepared by Dr. Fuxin Li, Oregon State Univ.

*With materials from Zolt Kira, Roger Grosse, Nitish Srivastava, Michael Nielsen*

# XOR problem and linear classifier


- 4 points:  $X = [(-1,-1), (-1,1), (1,-1), (1,1)]$
- $Y = [-1 \ 1 \ 1 \ -1]$
- Try using binomial log-likelihood loss:

$$\min_{\mathbf{w}, b} \sum_{i=1}^n \log(1 + e^{2y_i(\mathbf{w}^\top \mathbf{x}_i + b)})$$

- Gradient:

$$\nabla_{\mathbf{w}} = \sum_{i=1}^n 2y_i e^{2y_i(\mathbf{w}^\top \mathbf{x}_i + b)} / (1 + e^{2y_i(\mathbf{w}^\top \mathbf{x}_i + b)})$$

$$\nabla_b = \sum_{i=1}^n 2y_i e^{2y_i(\mathbf{w}^\top \mathbf{x}_i + b)} / (1 + e^{2y_i(\mathbf{w}^\top \mathbf{x}_i + b)})$$



Try  $\mathbf{w}=0, b=0$ , what

$$\nabla b = \sum_{i=1}^n 2y_i e^{2y_i(\mathbf{w}^\top \mathbf{x}_i + b)} / (1 + e^{2y_i(\mathbf{w}^\top \mathbf{x}_i + b)})$$


Try  $\mathbf{w}=0, b=0$ , what

# With 1 hidden layer


- A hidden layer makes a nonlinear classifier

$$f(x) = \mathbf{w}^\top g(\mathbf{W}^\top \mathbf{x} + \mathbf{c}) + b$$

- $g$  needs to be nonlinear
- Sigmoid:  $\text{Sigm}(x) = 1/(1 + e^{-x})$



- RELU:  $g(x) = \max(0, x)$



# Taking gradient

$$\min_{\mathbf{W}, \mathbf{w}} E(f) = \sum_i L(f(\mathbf{x}_i), y_i)$$

- What is  $\partial E / \partial \mathbf{W}$ ?  $f(x) = \mathbf{w}^\top g(\mathbf{W}^\top \mathbf{x} + \mathbf{c}) + b$
- Consider chain rule:  $dz/dx = dz/dy \, dy/dx$

# Note: Vectorized Computations

On the left are the computations performed by a network. Write them in terms of matrix and vector operations. Let  $\sigma(\mathbf{v})$  denote the logistic sigmoid function applied elementwise to a vector  $\mathbf{v}$ . Let  $\mathbf{W}$  be a matrix where the  $(i, j)$  entry is the weight from visible unit  $j$  to hidden unit  $i$ .

$$z_i = \sum_j w_{ij} x_j$$

$$h_i = \sigma(z_i)$$

$$y = \sum_i v_i h_i$$

$$\mathbf{z} = \mathbf{W}\mathbf{x}$$

$$\mathbf{h} = \sigma(\mathbf{z})$$

$$y = \mathbf{v}^T \mathbf{h}$$


# Backpropagation

- Save the gradients and the gradient products that have already been computed to avoid computing multiple times
- In a multiple layer network:
  - (Ignore constant terms)

$$f(x) = \mathbf{w}^{\top} \mathbf{n} \uparrow \tau g(\mathbf{W}^{\top} \mathbf{n} - \mathbf{1} \uparrow \tau g(\mathbf{W}^{\top} \mathbf{n} - \mathbf{2} \uparrow \tau g(\dots (\mathbf{W}^{\top} \mathbf{1} \uparrow \tau g(\mathbf{x}))))))$$

$$\begin{aligned} \partial E / \partial \mathbf{W}^{\top} k &= \partial E / \partial f \partial f / \partial f^{\top} k g(f^{\top} k - 1 (x)) \\ &= \partial E / \partial f^{\top} k + 1 \partial f^{\top} k + 1 / \partial f^{\top} k g(f^{\top} k - 1 (x)) \end{aligned}$$


$$f^{\top} k (x) = \mathbf{w}^{\top} k \uparrow \tau g(f^{\top} k - 1 (x)), f^{\top} 0 (x) = x$$



# Modules

- Each layer can be seen as a module
- Given input, return
  - Output  $f_{\downarrow a}(x)$
  - Network gradient  $\partial f_{\downarrow a} / \partial x$
  - Gradient of module parameters  $\partial f_{\downarrow a} / \partial u$
- During backprop, propagate/update
  - Backpropagated gradient

$$\partial E / \partial f_{\downarrow a}$$

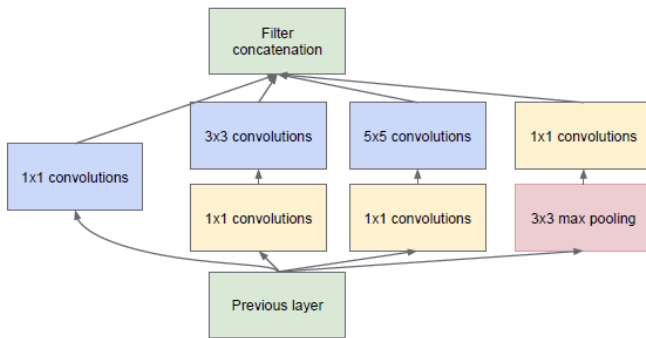


$$\begin{aligned} \partial E / \partial \mathbf{W}_{\downarrow k} &= \partial E / \partial f \partial f / \partial f_{\downarrow k} g(f_{\downarrow k-1}(x)) \\ &= \partial E / \partial f_{\downarrow k+1} \partial f_{\downarrow k+1} / \partial f_{\downarrow k} g(f_{\downarrow k-1}(x)) \end{aligned}$$

# Different DAG structures


- The backpropagation algorithm would work for any DAGs
- So one can imagine different architectures than the plain layerwise one

## Inception




(b) Inception module with dimension reductions

## Residual



## RNN






# Loss functions

- Regression:

- Least squares  $L(f) = (f(x) - y)^2$
- L1 loss  $L(f) = |f(x) - y|$
- Huber loss  $L(f) = \begin{cases} \frac{1}{2} (f(x) - y)^2, & |f(x) - y| < \delta \\ \delta |f(x) - y|, & \text{otherwise} \end{cases}$

- Binary Classification

- Hinge loss  $L(f) = \max(1 - yf(x), 0)$
- Binomial log-likelihood  $L(f) = \ln(1 + \exp(-2yf(x)))$
- Cross-entropy  $L(f) = y^* \ln(\text{sigm}(f)) + (1 - y^*) \ln(1 - \text{sigm}(f))$ ,
  - $y^* = (y + 1)/2$



# Multi-class: Softmax layer

- Multi-class logistic loss function

$$P(y = j | \mathbf{x}) = \frac{e^{\mathbf{x}^\top \mathbf{w}_j}}{\sum_{k=1}^K e^{\mathbf{x}^\top \mathbf{w}_k}}$$

- Log-likelihood:

- Loss function is minus log-likelihood

$$-\log P(y=j|x) = -\mathbf{x}^\top \mathbf{w}_j + \log \sum_{k=1}^K e^{\mathbf{x}^\top \mathbf{w}_k}$$


# Subgradients

- What if the function is non-differentiable?
- Subgradients:

- For **convex**  $f(x)$  at  $x \neq 0$  :
- If for any  $x$ ,


$$f(y) \geq f(x) + g^\top (y - x)$$

- $g$  is called a subgradient
- Subdifferential:  $\partial f$ : set of all subgradients
- Optimality condition:  $0 \in \partial f$



# The RELU unit

- $f(x) = \max(x, 0)$
- Convex
- Non-differentiable
- Subgradient:  $df/dx = \begin{cases} 1, & x > 0 \\ 0, & x \leq 0 \end{cases}$



# Subgradient descent

- Similar to gradient descent

$$x^{(k+1)} = x^{(k)} - \alpha_k g^{(k)}$$

- Step size rules:

- Constant step size:  $\alpha_k = \alpha$ .

- Square summable:  $\alpha_k \geq 0, \quad \sum_{k=1}^{\infty} \alpha_k^2 < \infty, \quad \sum_{k=1}^{\infty} \alpha_k = \infty.$

- Usually, a large constant that drops slowly after a long while
  - e.g.  $100/100+k$

# Universal Approximation Theorems

- Many universal approximation theorems proved in the 90s
- Simple statement: for every continuous function, there exist a function that can be approximated by a 1-hidden layer neural network with arbitrarily high precision

Formal statement [\[edit\]](#)

---

The theorem<sup>[\[2\]](#)[\[3\]](#)[\[4\]](#)[\[5\]](#)</sup> in mathematical terms:

Let  $\varphi(\cdot)$  be a nonconstant, [bounded](#), and [monotonically](#)-increasing [continuous](#) function. Let  $I_m$  denote the  $m$ -dimensional [unit hypercube](#)  $[0, 1]^m$ . The space of continuous functions on  $I_m$  is denoted by  $C(I_m)$ . Then, given any function  $f \in C(I_m)$  and  $\varepsilon > 0$ , there exists an integer  $N$  and real constants  $v_i, b_i \in \mathbb{R}$ , where  $i = 1, \dots, N$  such that we may define:

$$F(x) = \sum_{i=1}^N v_i \varphi(w_i^T x + b_i)$$

as an approximate realization of the function  $f$  where  $f$  is independent of  $\varphi$ ; that is,

$$|F(x) - f(x)| < \varepsilon$$

for all  $x \in I_m$ . In other words, functions of the form  $F(x)$  are [dense](#) in  $C(I_m)$ .

It obviously holds replacing  $I_m$  with any compact subset of  $\mathbb{R}^m$ .

# Universal Approximation Theorems

- The approximation does not need many units if the function is kinda nice. Let

$$C(f) = \int_{\mathbf{R}^d} |\omega| |f(\omega)| d\omega$$

- Then for a 1-hidden layer neural network with  $n$  hidden nodes, we have for a finite ball with radius  $r$ ,

$$\int_{B(r)} (f(x) - f_n(x))^2 d\mu(x) \leq 4r^2 C(f)^2 / n$$