
Agents and Environments

Lecture 2

Jim Rehg

January 16, 2016

College of Computing
Georgia Tech

(Slides based in part on slides from Andrea Thomaz)

Remaining Administrative Items

Resources

Review of Syllabus

Resources

Syllabus is now on T-Square, Piazza is coming soon

Text:

“AI: A Modern Approach,” 3rd edition (2010)
by Stuart Russell (UC Berkeley) & Peter Norvig (Google)

We will follow it reasonably closely in the beginning

Problems will be assigned from it

You must buy it

Grading

Midterm (30%)

Final (35%)

Projects (35%)

- 4 Graded Projects


- 1 Ungraded Project (Python Tutorial)

All projects done in Python (autograder)

No credit given for late submissions

Turn in on-time for Partial Credit

Honor Code: Projects are not collaborative, we test for cheating



Extra Credit

There will be opportunities for extra credit

Extra credit questions on all projects

Opportunities to participate in research studies

You must write a report

Don't rely on this!

Start Early Start Early Start Early Start Early

Schedule and Disclaimer

A complete lecture schedule will be on T-Square by Friday

This schedule is subject to change

Any changes announced in class & via T-Square

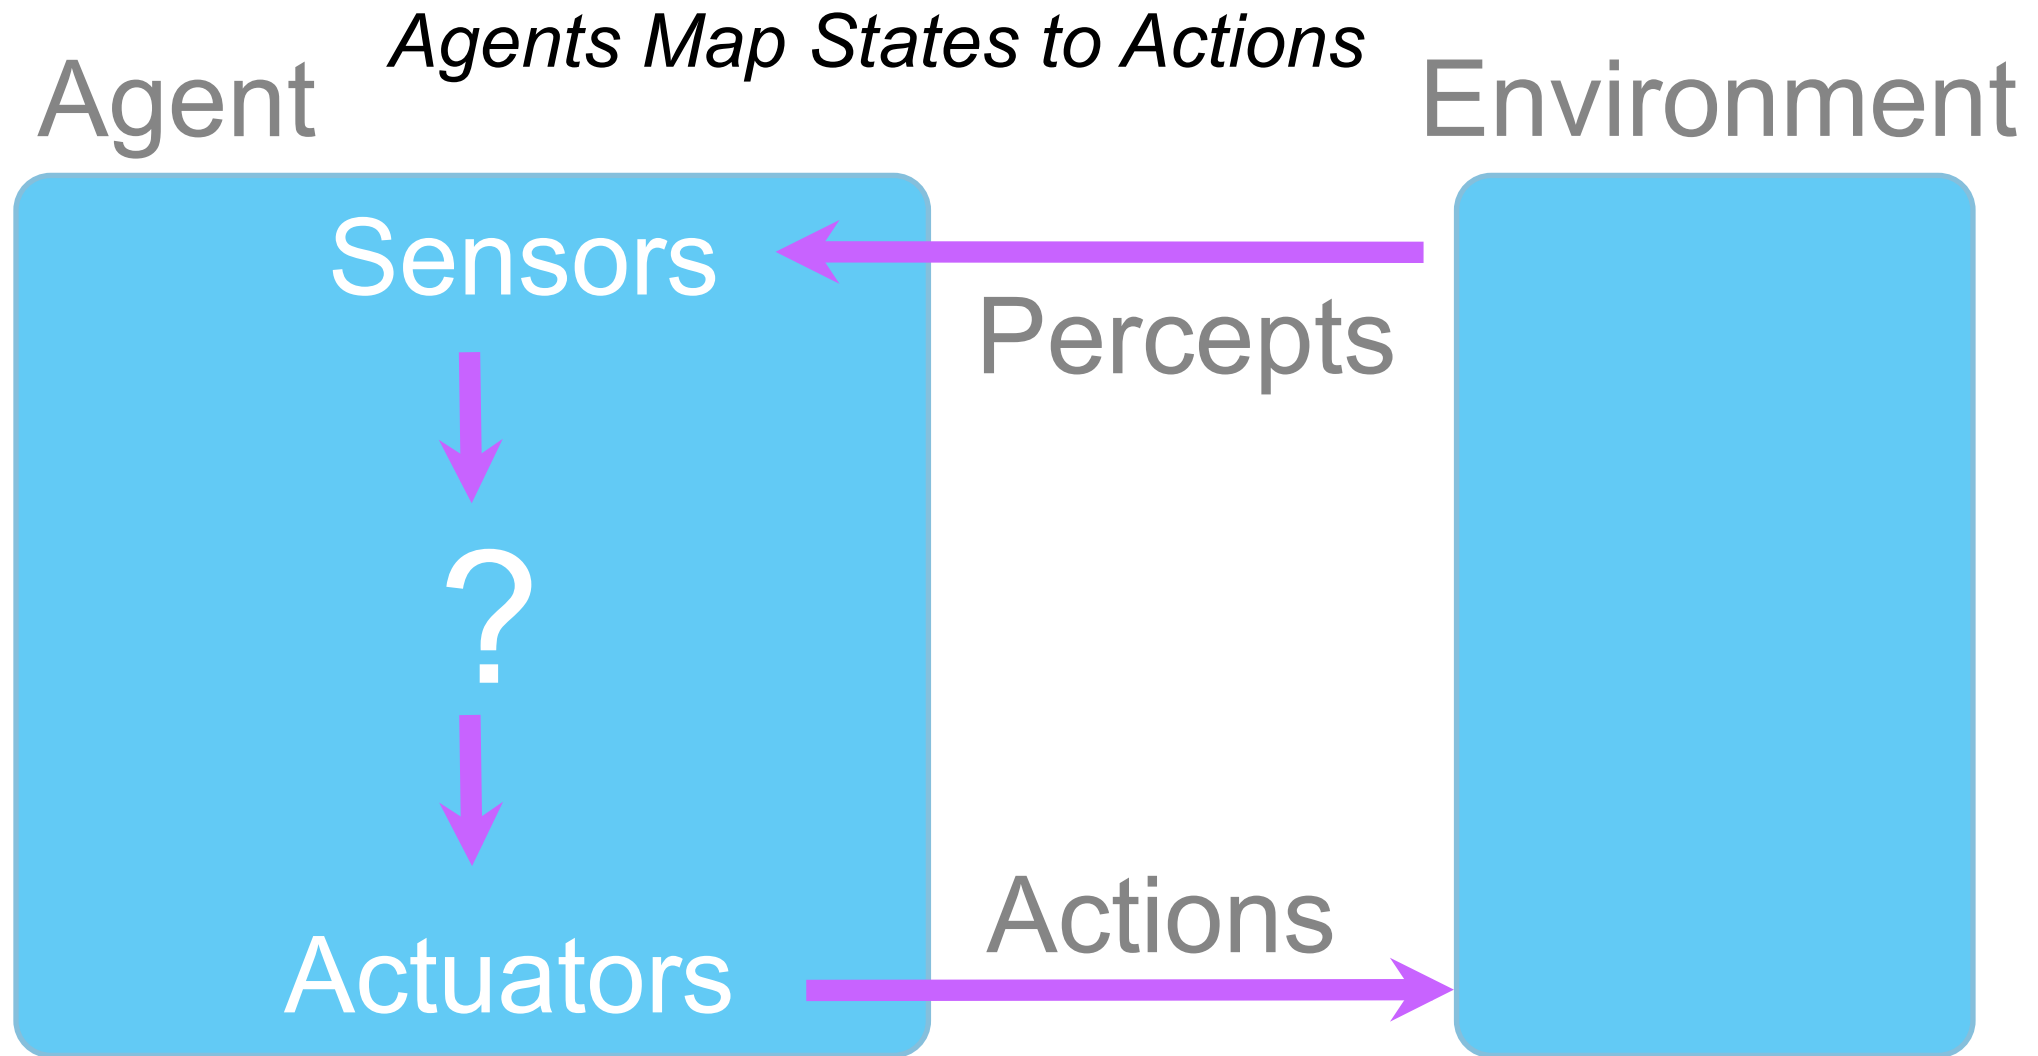
I will miss some classes due to unavoidable travel

I will tell you where I am going and why ☺

My replacement lecturer will be highly trained (usually faculty)


I will provide pre-recorded lectures in every case

Agents as Functions




Tabletop Robot Example

State (perception) <Ball-loc, Sq-loc, Hand-loc>	Action <L, R, Pick(obj), Drop>
Ball-A, Sq-Hand, Hand-A	




Tabletop Robot Example

State (perception) <Ball-loc, Sq-loc, Hand-loc>	Action <L, R, Pick(obj), Drop>
Ball-A, Sq-Hand, Hand-A	R
Ball-A, Sq-Hand, Hand-B	
	I



Tabletop Robot Example

State (perception) <Ball-loc, Sq-loc, Hand-loc>	Action <L, R, Pick(obj), Drop>
Ball-A, Sq-Hand, Hand-A	R
Ball-A, Sq-Hand, Hand-B	Drop
Ball-A, Sq-B, Hand-B	



Rational Agent

Always chooses actions to get the Best Expected Outcome



How desirable is this action?

Need a task-dependent measure of performance

For a state sequence S , select action A which maximizes the performance measure

Task Setting

The problem the agent is trying to solve:

P – Performance Metric

E – Environment

A – Actuators

S – Sensors

Automated Taxi Problem (coming soon...)

P – ?

E – ?

A – ?

S – ?

Automated Taxi Problem (coming soon...)

P – Safe, fast, legal, max revenue, min cost, min fuel, ...

E – city roads, traffic, pedestrians, bikers, construction, ...

A – Car controls (steering, gas pedal) and human interface

S – Cameras, radar, laser rangefinder, GPS, mapping, engine sensors, human input devices

Environment Types

Fully Observable

Deterministic

Episodic

Static

Discrete

Single-Agent

Partially Observable

Stochastic

Sequential

Dynamic

Continuous

Multi-Agent

Summary

Agents map states to actions

States and actions represented as tuples

Agents should be rational

Select the action that maximizes the outcome

Task defined by PEAS

Performance metric, Environment, Actuation, Sensors

Environment Types influence the difficulty of the problem

Questions?
