

Making Complex Decisions

Markov Decision Processes
CH 17

Section Overview

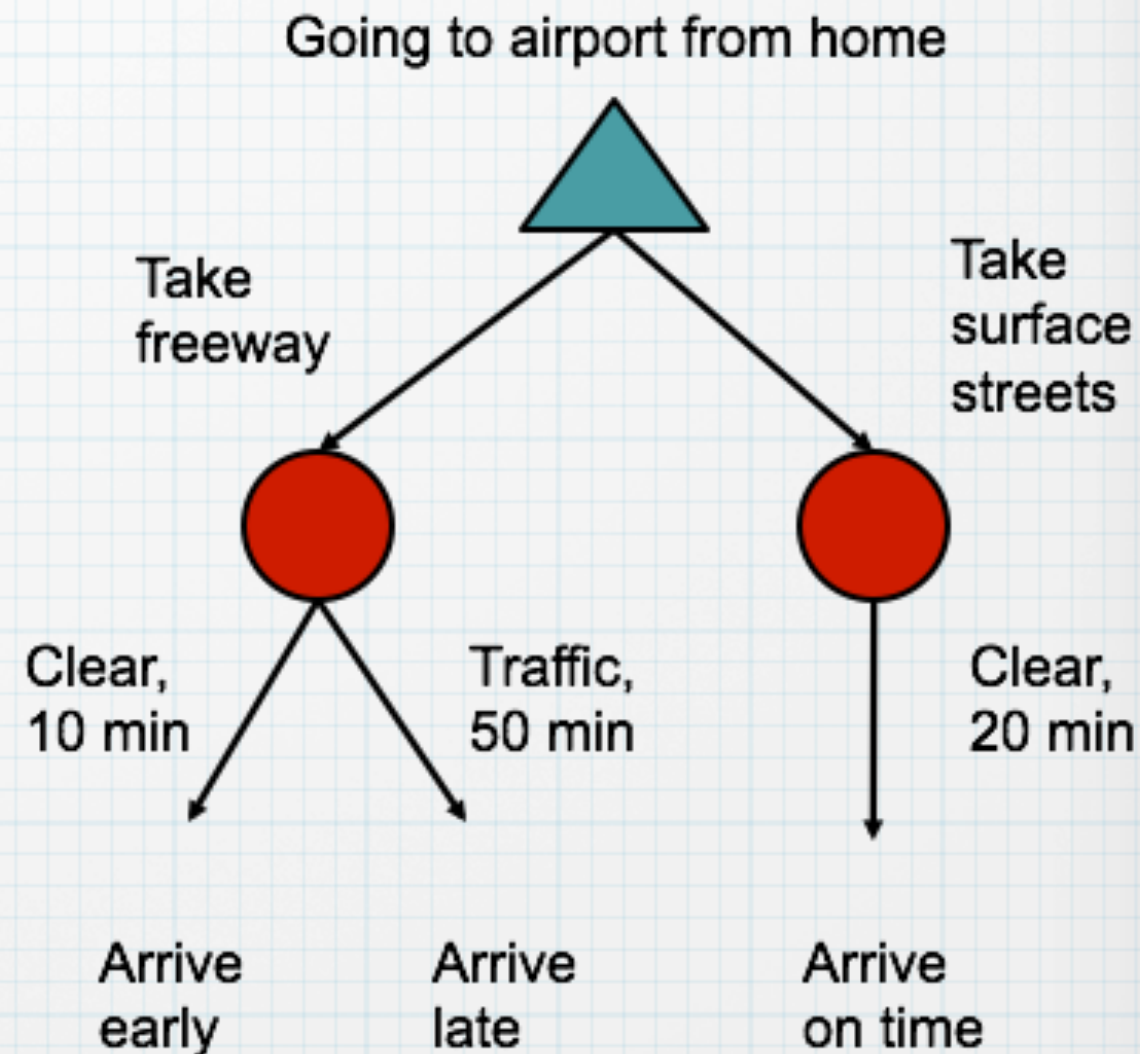
- * **Representing preferences**
- * **Markov Decision Processes**
- * **Solving MDPs**
 - * **Value Iteration**
 - * **Policy Iteration**
- * **Partially Observable MDPs**

Maximum Expected Utility

- * **Rational Agent = Maximizes its expected utility given its knowledge**
- * **Questions:**
 - * **Where do utilities come from?**
 - * **Why expected utility?**

Example

- * One way has a chance to be better or worse
- * How to decide?
- * Which would you pick if you are catching a flight?
- * Which if you are picking up a friend?



Assigning relative value to outcomes = Utilities

Agent Rational Decisions



- * Representing Decisions, Maximize Utility
 - * Receive feedback = rewards
 - * Utility defined by the reward function
 - * Act to maximize expected rewards
 - * Can learn to maximize rewards via Reinforcement Learning

Reward Functions

- * For example:
 - * Playing a game, reward defined at the end for winning or losing
 - * Vacuuming agent, reward for each piece of dirt
 - * Autonomous taxi, reward for each passenger delivered

MDP Grid World



- * Example agent for our MDP discussion
- * Walls block the agent
- * Actions only work 80% of the time
- * Big rewards at the end



Markov Decision Processes

- * MDP defined by
 - * States $s \in S$
 - * Actions $a \in A$
 - * Transition function $T(s, a, s')$
 - * Reward function $R(s, a, s')$

Just like our old search formulation
but with non-det actions and
rewards



What makes it Markov?

- * Markov means: given the present state future and past are independent
- * Specifically for MDPs Markov means

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0)$$

=

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$


Solving MDPs

- * In deterministic single-agent search problem we solved for an optimal plan, sequence of actions
- * Now an optimal policy $\pi^* = S \rightarrow A$
 - * policy, π , gives an action for every state
 - * optimal policy, π^* , maximizes expected utility
 - * defines a reflex agent


Solving MDPs

- * In deterministic single-agent search problem we solved for an optimal plan, sequence of actions
- * Now an optimal policy $\pi^* = S \rightarrow A$


Optimal policy when
 $R(s, a, s') = -0.03$ for
all non-terminals s




Example Optimal Policies




$$R(s) = -0.01$$



$$R(s) = -0.03$$



$$R(s) = -0.4$$



$$R(s) = -2.0$$

Summary + Preview



- * Preferences & Utilities & Rewards
- * Markov Decision Processes
- * Next...
 - * Two algorithms for solving an MDP
 - * Partially observable states

Making Decisions

- * When agent will need to act in the same environment over and over, to achieve the same goal, with nondeterministic actions**
- * “Policy” of action instead of “Plans”**

Markov Decision Processes

- * MDP defined by
 - * States $s \in S$
 - * Actions $a \in A$
 - * Transition function $T(s, a, s')$
 - * Reward function $R(s, a, s')$





Trashcan Robot MDP



You are building a robot to move around an office building collecting trash (from tables, floors, people, etc.). Define how you would set up this task as a Markov Decision Process




- * What is the state space?
- * What are the actions needed?
- * What should the rewards be?


MDP Practice Exercise

- * What are the possible states the agent could be in after moving: $a_1 = \text{North}$, $a_2 = \text{North}$
- * With what probabilities?













$$(1,3) = .8 * .8 = .64$$








$$(1,3) = .8 * .8 = .64$$


$$(1,2) = .8 * .2 + .1 * .8 = .24$$








$$\begin{aligned}
 (1,3) &= .8 * .8 = .64 \\
 (1,2) &= .8 * .2 + .1 * .8 = .24 \\
 (1,1) &= .1 * .1 + .1 * .1 = .02
 \end{aligned}$$





$$\begin{aligned}
 (1,3) &= .8 * .8 = .64 \\
 (1,2) &= .8 * .2 + .1 * .8 = .24 \\
 (1,1) &= .1 * .1 + .1 * .1 = .02 \\
 (2,1) &= .1 * .1 + .1 * .8 = .09
 \end{aligned}$$





$P(s_2)$



$$(1,3) = .8 * .8 = .64$$

$$(1,2) = .8 * .2 + .1 * .8 = .24$$

$$(1,1) = .1 * .1 + .1 * .1 = .02$$

$$(2,1) = .1 * .1 + .1 * .8 = .09$$

$$(3,1) = .1 * .1 = .01$$




Markov Decision Process (MDP)

MDP Solution = policy: what action to do in every state

Expected Utility: the expected reward from executing a particular policy

Optimal Policy: has the highest expected utility

Utility of Sequences




- * **Finite**

- * fixed time the agent is around
- * the right action in a state depends on how much time left


- * **Infinite**

- * no time limits, optimal action only depends on the state



$$R(s) = -0.03$$

Utility of Sequences



* Finite

- * fixed time the agent is around
- * the right action in a state depends on how much time left

* Infinite

- * no time limits, optimal action only depends on the state

Non-stationary

Stationary

Utilities of Sequences

- * Problem: infinite sequences have infinite rewards

$$U([r_0, r_1, r_2, \dots]) = r_0 + r_1 + r_2 + \dots$$

- * Solution?
 - * cut-off sequences, to make finite
 - * ...non-stationary policies

Utilities of Sequences

- * Problem: infinite sequences have infinite rewards


$$U([r_0, r_1, r_2, \dots]) = r_0 + r_1 + r_2 + \dots$$

- * Solution: discount future rewards

$$U([r_0, r_1, r_2, \dots]) = r_0 + \gamma r_1 + \gamma^2 r_2 \dots$$


Discount Future

- * Discount factor for each time step $\gamma < 1$
- * Earlier rewards have higher utility than later rewards
- * With discounted rewards, the utility of an infinite sequence is finite



Optimal Values for States

- * Optimal values define optimal policies
- * $U^*(s)$ = expected utility starting in s and acting optimally
- * $Q^*(s,a)$ = expected utility taking a in s , and then acting optimally



3	0.812	0.868	0.912	$+1$
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4

3	→	→	→	$+1$
2	↑		↑	-1
1	↑	←	←	←
	1	2	3	4

Calculating Policies

MDPs end up being a good representation for many real world problems.

Given an MDP description, several algorithms for solving for the optimal policy

Two general classes of algorithms:

Value Iteration and Policy Iteration

Optimal Utility

- * **Utility or Value of a state**
 - * **Expected utility of the sequence to follow that state, with particular policy of action**
 - * **$U^*(s)$ is the expected utility of following the optimal policy from s**
- * **$U(s)$ vs. $R(s)$**

Optimal Policy

$$U^{\pi}(S) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t) \right]$$

$$\pi_S^* = \arg \max_{\pi} U^{\pi}(S)$$

$$\pi_{s_t}^* = \arg \max_{a \in A(s_t)} \sum_{s_{t+1}} P(s_{t+1} \mid s_t, a) U(s_{t+1})$$

Expected utility