Making Complex Decisions

Markov Decision Processes
CH 17

Section Overview

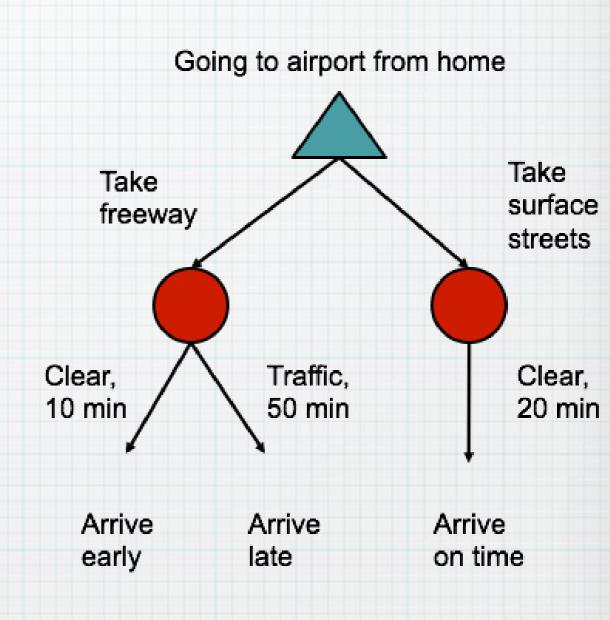
- Representing preferences
- * Markov Decision Processes
- Solving MDPs
 - * Value Iteration
 - Policy Iteration
- Partially Observable MDPs

Maximum Expected Utility

- Rational Agent = Maximizes its expected utility given its knowledge
- * Questions:
 - * Where do utilities come from?
 - * Why expected utility?

Example

- One way has a chance to be better or worse
- * How to decide?
- * Which would you pick if you are catching a flight?
- * Which if you are picking up a friend?



Assigning relative value to outcomes = Utilities

Agent Rational Decisions

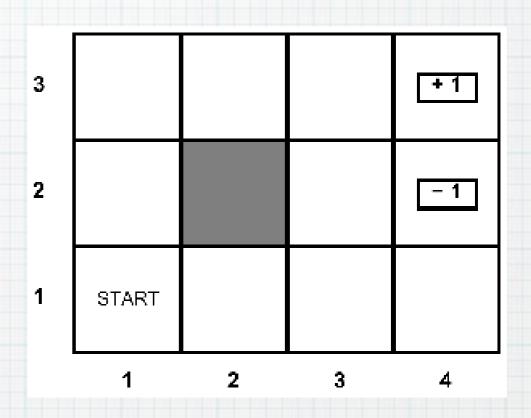
- Representing Decisions, Maximize
 Utility
 - * Receive feedback = rewards
 - Utility defined by the reward function
 - * Act to maximize expected rewards
 - Can learn to maximize rewards via Reinforcement Learning

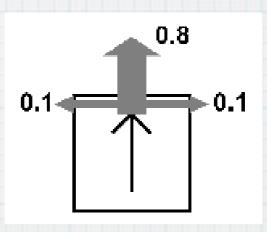
Reward Functions

- * For example:
 - Playing a game, reward defined at the end for winning or losing
 - Vacuuming agent, reward for each piece of dirt
 - Autonomous taxi, reward for each passenger delivered

MDP Grid World

- Example agent for our MDP discussion
- Walls block the agent
- Actions only work 80% of the time
- Big rewards at the end

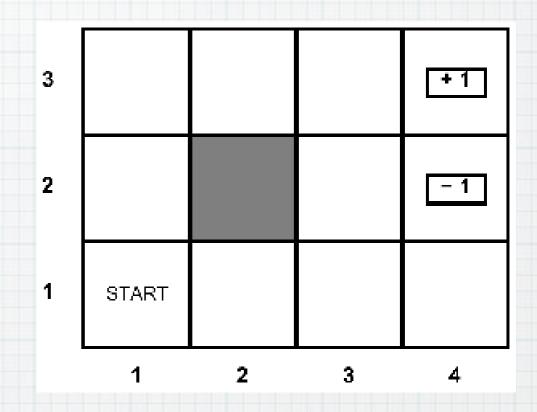


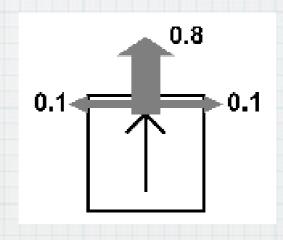


Markov Decision Processes

- MDP defined by
 - ***** States s∈S
 - * Actions a∈A
 - Transition functionT(s,a,s')
 - Reward function R(s,a,s')

Just like our old search formulation but with non-det actions and rewards





What makes it Narkov?

- Markov means: given the present state future and past are independent
- Specifically for MDPs Markov means

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots S_0 = s_0)$$

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$

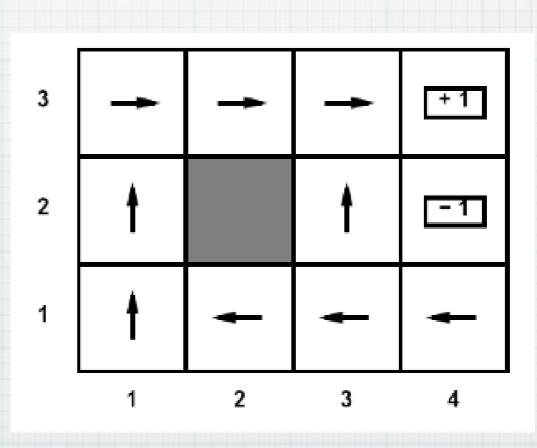
Solving MDPs

- In deterministic single-agent search problem we solved for an optimal plan, sequence of actions
- * Now an optimal policy $\Pi^* = S \rightarrow A$
 - * policy, \(\preceq\), gives an action for every state
 - optimal policy, \(\pi\)*, maximizes expected utility
 - * defines a reflex agent

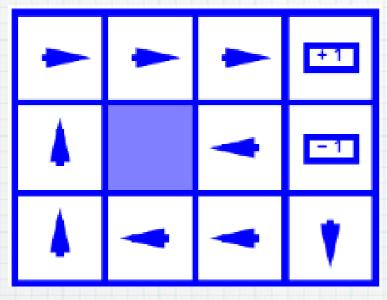
Solving MDPs

- In deterministic single-agent search problem we solved for an optimal plan, sequence of actions
- * Now an optimal policy $\Pi^* = S \rightarrow A$

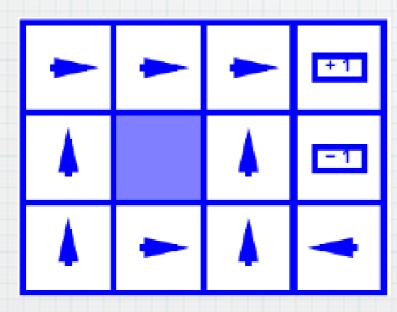
Optimal policy when R(s, a, s') = -0.03 for all non-terminals s



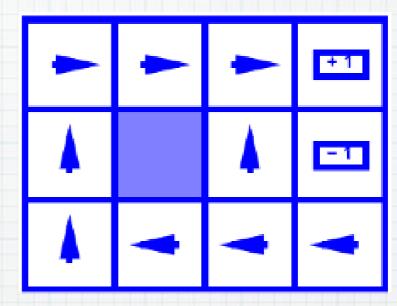
Example Optimal Policies



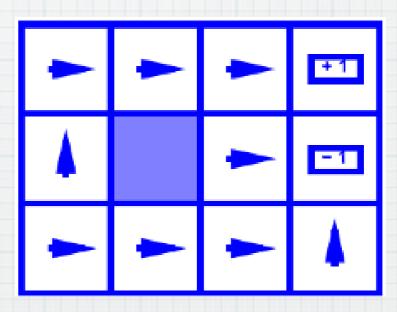
R(s) = -0.01



R(s) = -0.4



R(s) = -0.03



$$R(s) = -2.0$$

Summary + Preview

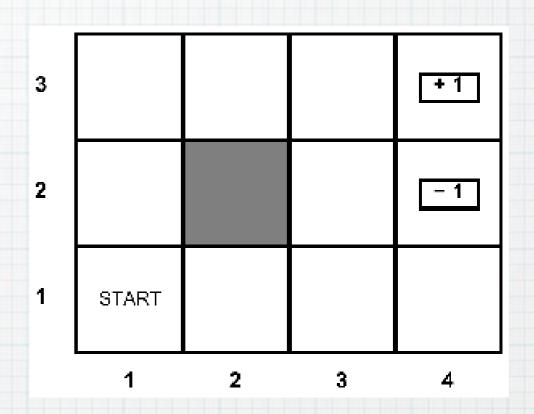
- * Preferences & Utilities & Rewards
- * Markov Decision Processes
- * Next...
 - * Two algorithms for solving an MDP
 - Partially observable states

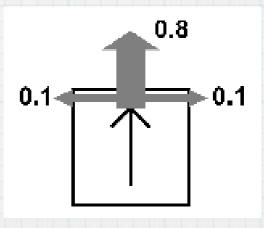
Making Decisions

- * When agent will need to act in the same environment over and over, to achieve the same goal, with nondeterministic actions
- * "Policy" of action instead of "Plans"

Markov Decision Processes

- MDP defined by
 - ∗ States s∈S
 - ⋆ Actions a∈A
 - Transitionfunction T(s,a,s')
 - Reward functionR(s,a,s')





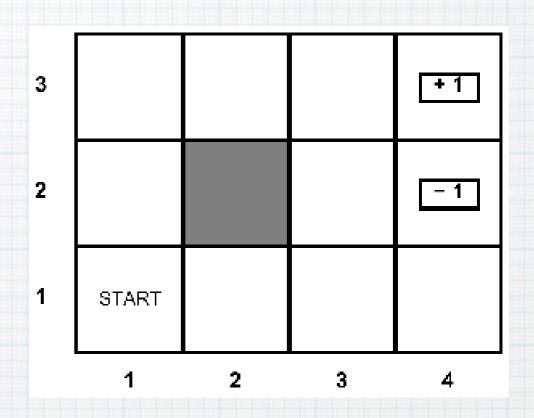
Trashcan Robot MDP

You are building a robot to move around an office building collecting trash (from tables, floors, people, etc.). Define how you would set up this task as a Markov Decision Process

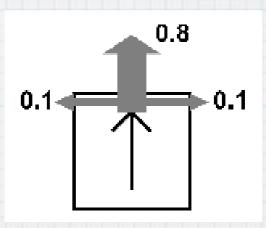
- * What is the state space?
- * What are the actions needed?
- * What should the rewards be?

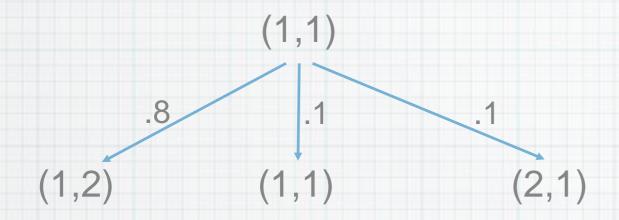
MDP Practice Exercise

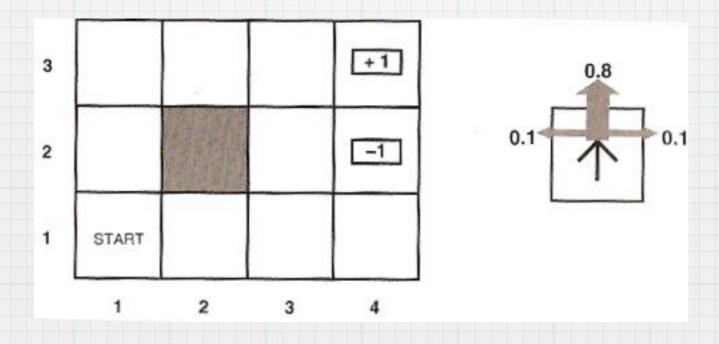
What are the possible states the agent could be in after moving:
 a₁=North, a₂=North

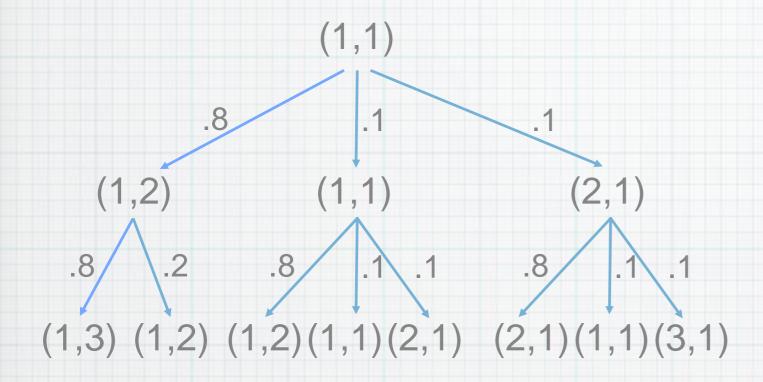


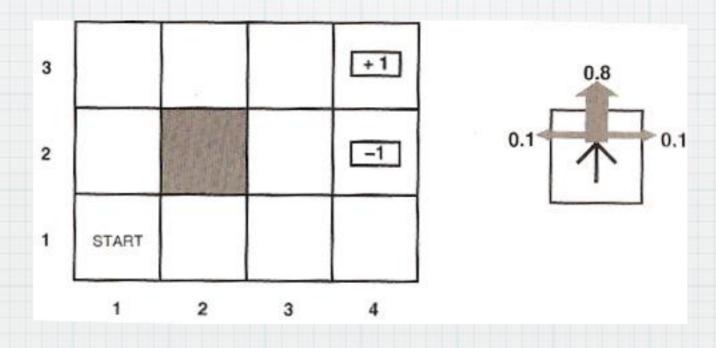
* With what probabilities?

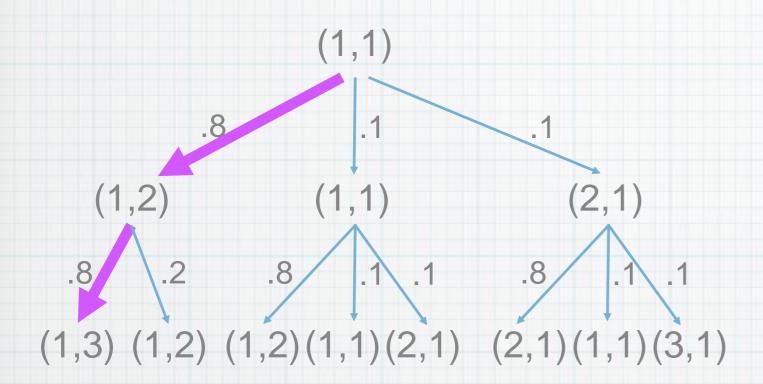




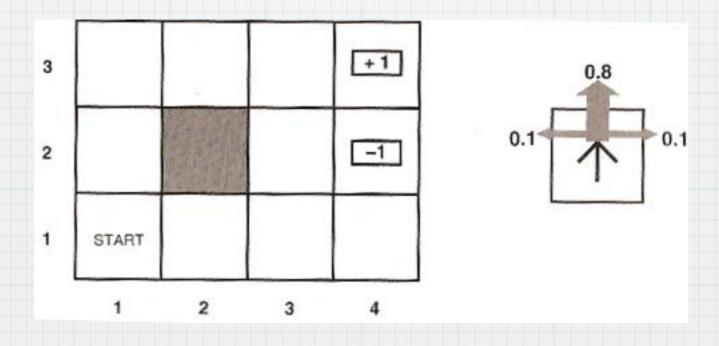


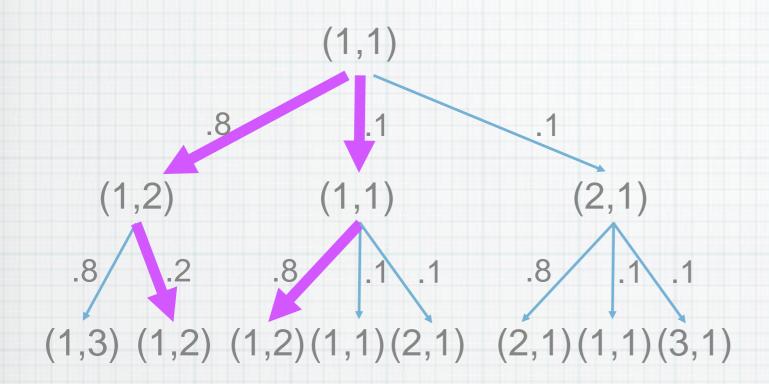






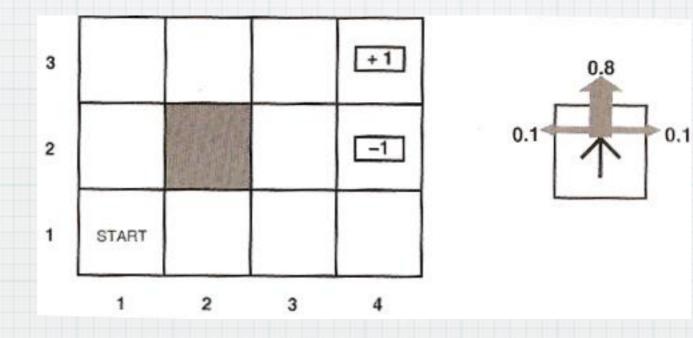
$$(1,3) = .8*.8 = .64$$

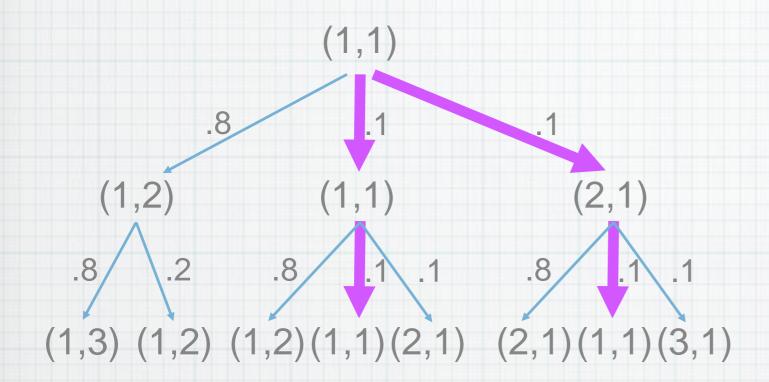




$$(1,3) = .8*.8 = .64$$

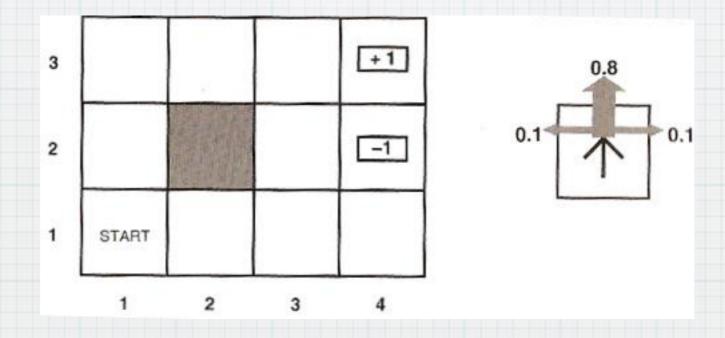
 $(1,2) = .8*.2 + .1*.8 = .24$

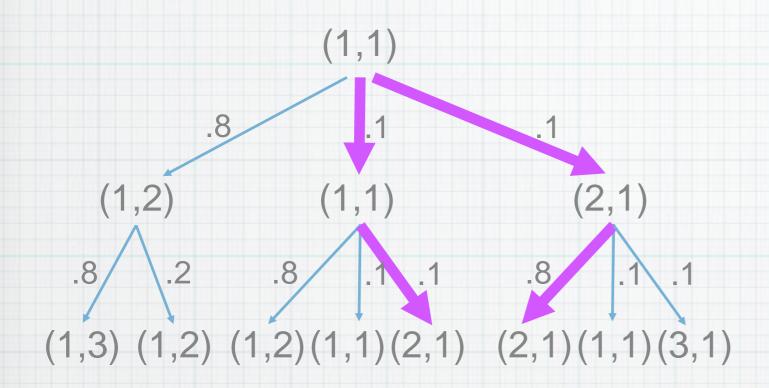




$$(1,3) = .8*.8 = .64$$

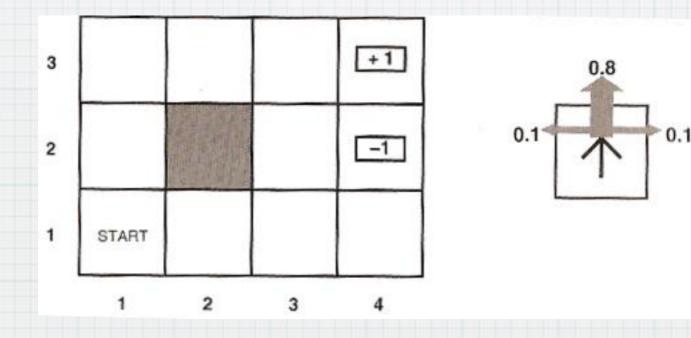
 $(1,2) = .8*.2 + .1*.8 = .24$
 $(1,1) = .1*.1 + .1*.1 = .02$

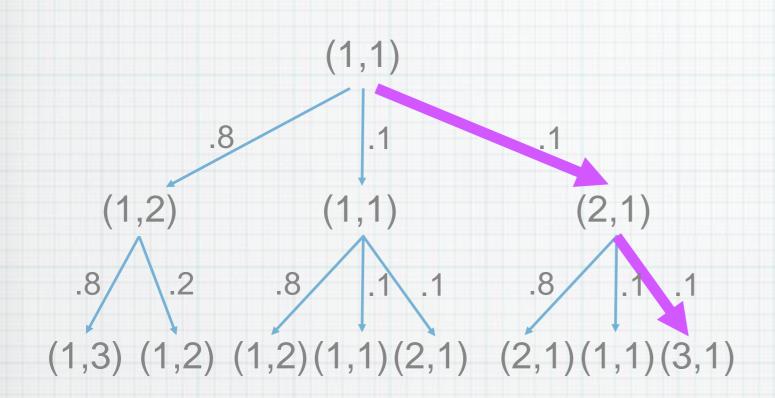




$$(1,3) = .8*.8 = .64$$

 $(1,2) = .8*.2 + .1*.8 = .24$
 $(1,1) = .1*.1 + .1*.1 = .02$
 $(2,1) = .1*.1 + .1*.8 = .09$





$P(s_2)$

$$(1,3) = .8*.8 = .64$$


$$(1,2) = .8*.2 + .1*.8 = .24$$

$$(1,1) = .1*.1 + .1*.1 = .02$$

$$(2,1) = .1*.1 + .1*.8 = .09$$

$$(3,1) = .1*.1 = .01$$

0.1



Markov Decision Process (MDP)

MDP Solution = policy: what action to do in every state

Expected Utility: the expected reward from executing a particular policy

Optimal Policy: has the highest expected utility

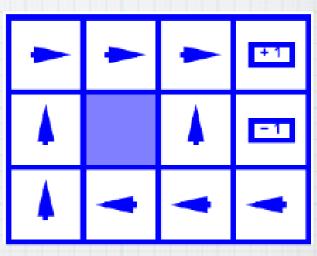
Utility of Sequences

* Finite

- * fixed time the agent is around
- * the right action in a state depends on how much time left

* Infinite

* no time limits, optimal action only depends on the state



R(s) = -0.03

Utility of Sequences

* Finite

- * fixed time the agent is around
- * the right action in a state depends on how much time left

* Infinite

Non-stationary

no time limits, optimal action only depends on the state

Stationary

Utilities of Sequences

* Problem: infinite sequences have infinite rewards

$$U([r_0, r_1, r_2, \ldots]) = r_0 + r_1 + r_2 + \cdots$$

- * Solution?
 - cut-off sequences, to make finite
 - * ...non-stationary policies

Utilities of Sequences

* Problem: infinite sequences have infinite rewards

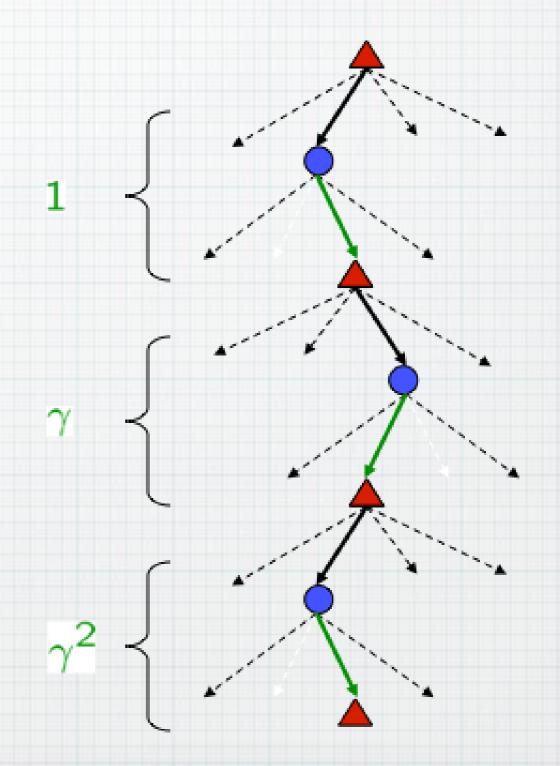
$$U([r_0, r_1, r_2, \ldots]) = r_0 + r_1 + r_2 + \cdots$$

Solution: discount future rewards

$$U([r_0, r_1, r_2, \ldots]) = r_0 + \gamma r_1 + \gamma^2 r_2 \cdots$$

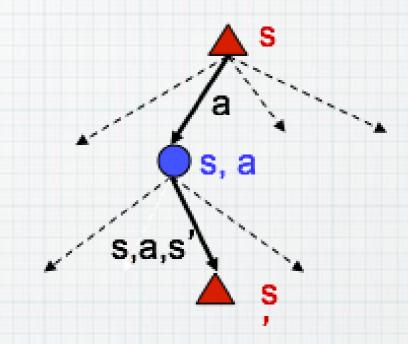
Discount Future

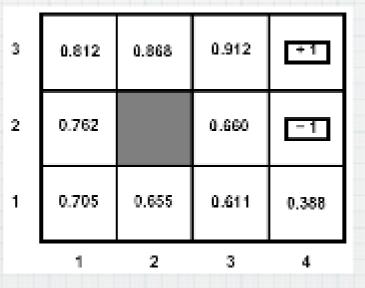
- Discount factor for each time step <<1
- Earlier rewards
 have higher utility
 than later rewards
- With discounted rewards, the utility of an infinite sequence is finite

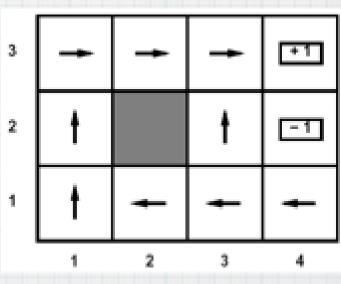


Optimal Values for States

- Optimal values define optimal policies
- U*(s) = expected
 utility starting in s
 and acting optimally
- * Q*(s,a) = expected utility taking a in s, and then acting optimally







Calculating Policies

MDPs end up being a good representation for many real world problems.

Given an MDP description, several algorithms for solving for the optimal policy

Two general classes of algorithms:

Value Iteration and Policy Iteration

Optimal Utility

- Utility or Value of a state
 - * Expected utility of the sequence to follow that state, with particular policy of action
 - U*(s) is the expected utility of following the optimal policy from s
- * U(s) vs. R(s)

Optimal Policy

$$U^{\pi}(S) = E\left[\sum_{t=0}^{\infty} \gamma^{t} R(S_{t})\right]$$

$$\pi_S^* = \arg\max_{\pi} U^{\pi}(S)$$

$$\pi_{s_t}^* = \underset{a \in A(s_t)}{\operatorname{arg\,max}} \sum_{s_{t+1}} P(s_{t+1} \mid s_t, a) U(s_{t+1})$$

Expected utility